

```
In [57]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression
```

```
In [58]: df = pd.read_csv('../Data/BostonHousing.csv')
df
```

```
Out[58]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7
...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0

506 rows × 14 columns

```
In [59]: df.isnull().sum()
```

```
Out[59]: crim      0
zn            0
indus         0
chas          0
nox           0
rm            0
age           0
dis           0
rad           0
tax           0
ptratio       0
b             0
lstat         0
medv          0
dtype: int64
```

```
In [65]: scale = MinMaxScaler()
Scaled_data = df.iloc[:, :-1] = scale.fit_transform(df.iloc[:, :-1])
df
```

Out [65]:

	crim	zn	indus	chas	nox	rm	age	dis
0	0.000000	0.18	0.067815	0	0.314815	0.577505	0.641607	0.269203
1	0.000236	0.00	0.242302	0	0.172840	0.547998	0.782698	0.348962
2	0.000236	0.00	0.242302	0	0.172840	0.694386	0.599382	0.348962
3	0.000293	0.00	0.063050	0	0.150206	0.658555	0.441813	0.448545
4	0.000705	0.00	0.063050	0	0.150206	0.687105	0.528321	0.448545
...	...	...	...	...	...	...	...	...
501	0.000633	0.00	0.420455	0	0.386831	0.580954	0.681771	0.122671
502	0.000438	0.00	0.420455	0	0.386831	0.490324	0.760041	0.105293
503	0.000612	0.00	0.420455	0	0.386831	0.654340	0.907312	0.094381
504	0.001161	0.00	0.420455	0	0.386831	0.619467	0.889804	0.114514
505	0.000462	0.00	0.420455	0	0.386831	0.473079	0.802266	0.125072

506 rows x 14 columns

In [61]:

```
plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True)
plt.show()
```



In [62]:

```
x = df.iloc[:, :-1]
y = df[["rm", "b", "zn"]]
```

```
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
model=LinearRegression()
model.fit(X_train,y_train)
y_test_pred = model.predict(X_test)
y_train_pred = model.predict(X_train)
comparison=pd.DataFrame({"Actual":y_test.values.flatten(),"Predicted":y_train_pred.values.flatten()})
comparison
```

Out [62]:

	Actual	Predicted
0	0.547040	5.470397e-01
1	0.996470	9.964698e-01
2	0.000000	4.926481e-16
3	0.612569	6.125695e-01
4	1.000000	1.000000e+00
...	...	...
301	0.997705	9.977054e-01
302	0.000000	8.702437e-16
303	0.522131	5.221307e-01
304	0.965530	9.655303e-01
305	0.000000	3.439613e-16

306 rows × 2 columns

In [63]:

```
rmse_test = np.sqrt(mean_squared_error(y_test, y_test_pred))
print(f"Test RMSE: {rmse_test}")
rmse_train = np.sqrt(mean_squared_error(y_train, y_train_pred))
print(f"Train RMSE: {rmse_train}")
```

Test RMSE: 7.71351067881442e-16

Train RMSE: 7.434339382369642e-16

In [64]:

```
comparison = pd.DataFrame({
    "Actual": y_test.values.flatten(),
    "Predicted": y_test_pred.flatten()
})

columns = ["rm", "b", "zn"]
train_colors = ['blue', 'green', 'orange']
test_colors = ['cyan', 'lightgreen', 'gold']

fig, axes = plt.subplots(1, 3, figsize=(18, 6))

for i, col in enumerate(columns):
    axes[i].scatter(y_train[col], y_train_pred[:, i], color=train_colors[i])
    axes[i].scatter(y_test[col], y_test_pred[:, i], color=test_colors[i])

    min_val = min(y[col].min(), y[col].min())
    max_val = max(y[col].max(), y[col].max())
    axes[i].plot([min_val, max_val], [min_val, max_val], color='black', linestyle='dashed')

    axes[i].set_xlabel("Actual Values")
    axes[i].set_ylabel("Predicted Values")
```

```
axes[i].set_title(f"Actual vs Predicted for {col}")
axes[i].legend()
axes[i].grid(True)
```

```
plt.tight_layout()
plt.show()
```

